

《操作系统原理》实验报告

姓名：韩艺轩
学院：计算机科学与技术学院
专业：计算机科学与技术（图灵班）
邮箱：2674361965@qq.com
指导教师：申文博
报告日期：2023/10/26

- 实验步骤
 - 准备工作
 - proc.h 数据结构定义
 - 线程调度功能实现
 - 线程初始化
 - __dummy 与 dummy
 - 实现线程切换
 - 实现调度入口函数
 - 实现线程调度
 - 短作业优先调度算法
 - 优先级调度算法
- 编译及测试
 - 测试 短作业优先调度算法
 - NR_TASKS = 4
 - NR_TASKS = 8
 - NR_TASKS = 16
 - 测试 优先级调度算法
 - NR_TASKS = 4
 - NR_TASKS = 8
 - NR_TASKS = 16
 - 运行 短作业优先调度算法
 - 运行 优先级调度算法
- 思考题

实验步骤

准备工作

从 repo 同步以下代码: rand.h/rand.c, string.h/string.c, mm.h/mm.c, proc.h/proc.c, test.h/test_schedule.h, schedule_null.c/schedule_test.c 以及新增的一些 Makefile 的变化。

```
.
├── arch
│   └── riscv
│       ├── include
│       │   ├── mm.h
│       │   └── proc.h
│       └── kernel
│           ├── mm.c
│           └── proc.c
├── include
│   ├── rand.h
│   ├── string.h
│   ├── test.h
│   └── schedule_test.h
├── test
│   ├── schedule_null.c
│   ├── schedule_test.c
│   └── Makefile
├── lib
│   ├── rand.c
│   └── string.c
└── Makefile
```

修改 defs.h, 在 defs.h 添加如下内容:

```
#define PHY_START 0x0000000080000000
#define PHY_SIZE 128 * 1024 * 1024 // 128MB, QEMU 默认内存大小
#define PHY_END (PHY_START + PHY_SIZE)

#define PGSIZE 0x1000 // 4KB
#define PGROUNDUP(addr) ((addr + PGSIZE - 1) & ~(PGSIZE - 1))
#define PGROUNDDOWN(addr) (addr & ~(PGSIZE - 1))
```

proc.h 数据结构定义

对 proc.h 数据结构进行定义

```

// arch/riscv/include/proc.h

#include "types.h"

// #define NR_TASKS (1 + 31) // 用于控制 最大线程数量 (idle 线程 + 31 内核线程)
#define NR_TASKS (1 + 3)
#define TASK_RUNNING 0 // 为了简化实验，所有的线程都只有一种状态

#define PRIORITY_MIN 1
#define PRIORITY_MAX 10

/* 用于记录 `线程` 的 `内核栈与用户栈指针` */
/* (lab2 中无需考虑，在这里引入是为了之后实验的使用) */
struct thread_info {
    uint64 kernel_sp;
    uint64 user_sp;
};

/* 线程状态段数据结构 */
struct thread_struct {
    uint64 ra;
    uint64 sp;
    uint64 s[12];
};

/* 线程数据结构 */
struct task_struct {
    struct thread_info thread_info;
    uint64 state; // 线程状态
    uint64 counter; // 运行剩余时间
    uint64 priority; // 运行优先级 1最低 10最高
    uint64 pid; // 线程id

    struct thread_struct thread;
};

/* 线程初始化 创建 NR_TASKS 个线程 */
void task_init();

/* 在时钟中断处理中被调用 用于判断是否需要调度 */
void do_timer();

/* 调度程序 选择出下一个运行的线程 */
void schedule();

/* 线程切换入口函数*/
void switch_to(struct task_struct* next);

/* dummy function: 一个循环程序，循环输出自己的 pid 以及一个自增的局部变量 */
void dummy();

```

线程调度功能实现

线程初始化

当我们的 OS run 起来的时候，其本身就是一个线程 idle。第一步我们要为 idle 设置 task_struct。并将 current, task[0] 都指向 idle。为了方便起见，我们将 task[1] ~ task[NR_TASKS - 1]，全部初始化，这里和 idle 设置的区别在于要为这些线程设置 thread_struct 中的 ra 和 sp。

完善 arch/riscv/kernel/proc.c

```
void task_init() {
    test_init(NR_TASKS);

    idle = (struct task_struct *)kalloc();
    idle->state = TASK_RUNNING;
    idle->counter = 0;
    idle->priority = 0;
    idle->pid = 0;
    current = idle;
    task[0] = idle;

    for(uint64 i = 1; i < NR_TASKS; i++){
        task[i] = (struct task_struct *)kalloc();
        task[i]->state = TASK_RUNNING;
        task[i]->counter = task_test_counter[i];
        task[i]->priority = task_test_priority[i];
        task[i]->pid = i;
        task[i]->thread.ra = (uint64)__dummy ;
        task[i]->thread.sp = (uint64)task[i] + PGSIZE ;
    }
    printk("...proc_init done!\n");
}
```

__dummy 与 dummy

task[1] ~ task[NR_TASKS - 1]都运行同一段代码 dummy() 我们在 proc.c 添加 dummy():

```
// arch/riscv/kernel/proc.c

void dummy() {
    schedule_test();
    uint64 MOD = 1000000007;
    uint64 auto_inc_local_var = 0;
    int last_counter = -1;
    while(1) {
        if ((last_counter == -1 || current->counter != last_counter) && current->counter > 0) {
            if(current->counter == 1){
                --(current->counter);    // forced the counter to be zero if this
                thread is going to be scheduled
            }
        }
    }
```

```

    } // in case that the new counter is also
    1, leading the information not printed.
    last_counter = current->counter;
    auto_inc_local_var = (auto_inc_local_var + 1) % MOD;
    printk("[PID = %d] is running. auto_inc_local_var = %d\n", current-
>pid, auto_inc_local_var);
    }
}
}

```

当我们创建一个新的线程，此时线程的栈为空，当这个线程被调度时，是没有上下文需要被恢复的，所以我们需要为线程 第一次调度 提供一个特殊的返回函数 `__dummy`

在 `entry.S` 添加 `__dummy`

- 在 `__dummy` 中将 `sepc` 设置为 `dummy()` 的地址，并使用 `sret` 从中断中返回。
- `__dummy` 与 `_traps` 的 `restore` 部分相比，其实就是省略了从栈上恢复上下文的过程（但是手动设置了 `sepc`）。

```

__dummy:
    la t0, dummy
    csrw sepc, t0
    sret

```

实现线程切换

判断下一个执行的线程 `next` 与当前的线程 `current` 是否为同一个线程，如果是同一个线程，则无需做任何处理，否则调用 `__switch_to` 进行线程切换。

```

// arch/riscv/kernel/proc.c

void switch_to(struct task_struct* next) {
    if((uint64)current != (uint64)next){
        struct task_struct *prev = current;
        current = next;
#ifdef SJF
        printk("\n");
        printk("switch to [PID = %d COUNTER = %d]\n\n", current->pid, current-
>counter);
#endif
#ifdef PRIORITY
        printk("\n");
        printk("switch to [PID = %d PRIORITY = %d COUNTER = %d]\n\n", current-
>pid, current->priority, current->counter);
#endif
        __switch_to(prev, next);
    }
}

```

在 entry.S 中实现线程上下文切换 __switch_to:

- __switch_to接受两个 task_struct 指针作为参数保存当前线程的ra, sp, s0~s11到当前线程的 thread_struct 中
- 将下一个线程的 thread_struct 中的相关数据载入到ra, sp, s0~s11中。

```
# arch/riscv/kernel/entry.S
```

```
__switch_to:
    sd ra, 48(a0)
    sd sp, 56(a0)
    sd s0, 64(a0)
    sd s1, 72(a0)
    sd s2, 80(a0)
    sd s3, 88(a0)
    sd s4, 96(a0)
    sd s5, 104(a0)
    sd s6, 112(a0)
    sd s7, 120(a0)
    sd s8, 128(a0)
    sd s9, 136(a0)
    sd s10, 144(a0)
    sd s11, 152(a0)

    ld ra, 48(a1)
    ld sp, 56(a1)
    ld s0, 64(a1)
    ld s1, 72(a1)
    ld s2, 80(a1)
    ld s3, 88(a1)
    ld s4, 96(a1)
    ld s5, 104(a1)
    ld s6, 112(a1)
    ld s7, 120(a1)
    ld s8, 128(a1)
    ld s9, 136(a1)
    ld s10, 144(a1)
    ld s11, 152(a1)

    ret
```

实现调度入口函数

实现 do_timer()

```
void do_timer(void) {
    // 1. 如果当前线程是 idle 线程 直接进行调度
    // 2. 如果当前线程不是 idle 对当前线程的运行剩余时间减1 若剩余时间仍然大于0 则直接返
```

回 否则进行调度

```

    if(current == idle){
        schedule();
    }else{
        if((signed long)(--(current->counter)) > 0 ){
            // printk("in do_timer COUNTER = %d\n",current->counter);
            return ;
        }else{
            schedule();
        }
    }
}

```

并在 时钟中断处理函数 中调用。

```

void trap_handler(unsigned long scause, unsigned long sepc) {
    unsigned long temp = 1;
    if((scause&(temp<<63))==(temp<<63) && scause&(temp<<4)==(temp<<4)){

        clock_set_next_event();

        do_timer();

        printk("[S] Supervisor Timer Interrupt.\n");
    }
}

```

实现线程调度

实现两种调度算法：1.短作业优先调度算法，2.优先级调度算法。

短作业优先调度算法

- 遍历线程指针数组task（不包括 idle，即 task[0]），在所有运行状态（TASK_RUNNING）下的线程运行剩余时间最少的线程作为下一个执行的线程。
- 如果所有运行状态下的线程运行剩余时间都为0，则对 task[1] ~ task[NR_TASKS-1] 的运行剩余时间重新赋值（使用 rand()），之后再重新进行调度。

```

// arch/riscv/kernel/proc.c
// void schedule(void)

#ifdef SJF
    uint64 judge = 1;
    for(uint64 i = 1;i<NR_TASKS;i++){
        if(task[i]&&task[i]->counter != 0){
            judge = 0;

```



```

    }
}
if(judge){
    // printk("in schedule() task[i]->counter are all 0!\n");
    for(uint64 i=1;i<NR_TASKS;i++){
        if(task[i]) task[i]->counter = rand();
    }
}
uint64 min = -1;
// printk("MIN == %d\n",min);
uint64 index = 0;
for(uint64 i = 1;i < NR_TASKS;i++){
    if(task[i]&&task[i]->counter < min && task[i]->state == TASK_RUNNING &&
task[i]->counter != 0){
        min = task[i]->counter;
        index = i;
    }
}
// printk("MIN ' == %d\n",min);
// printk("index = %d\n",index);
switch_to(task[index]);

#endif

```

优先级调度算法

本算法实现参考了 Linux v0.11 调度算法实现。

```

// arch/riscv/kernel/proc.c
// void schedule(void)

#ifdef PRIORITY
    uint64 max = task[1]->counter;
    uint64 index = 1;
    while(1){
        for(uint64 i=1;i<NR_TASKS;i++){
            if(task[i] && task[i]->counter > max && task[i]->state ==
TASK_RUNNING){
                max = task[i]->counter;
                index = i;
            }
        }
        if(max) break;
        for(uint64 i = 1;i<NR_TASKS;i++){
            if(task[i]){
                task[i] ->counter = (task[i]->counter >> 1) + task[i]->priority;
            }
        }
    }
    // printk("index == %d\n",index);

```

编译及测试

NR_TASKS = 4

NR_TASKS = 8

NR_TASKS = 16

测试 优先级调度算法

NR_TASKS = 4

```

C
CC[S] Supervisor Timer Interrupt.
C
CCC[S] Supervisor Timer Interrupt.
C
CCCC[S] Supervisor Timer Interrupt.
C
CCCCC[S] Supervisor Timer Interrupt.
C
CCCCCC[S] Supervisor Timer Interrupt.
C
CCCCCCC[S] Supervisor Timer Interrupt.
C
CCCCCCCC[S] Supervisor Timer Interrupt.
C
CCCCCCCCC
switch to [PID = 3 PRIORITY = 52 COUNTER = 8]

D
CCCCCCCCCD[S] Supervisor Timer Interrupt.
D
CCCCCCCCDD[S] Supervisor Timer Interrupt.
D
CCCCCCCCDDDD[S] Supervisor Timer Interrupt.
D
CCCCCCCCDDDD[S] Supervisor Timer Interrupt.
D
CCCCCCCCDDDDDD[S] Supervisor Timer Interrupt.
D
CCCCCCCCDDDDDD[S] Supervisor Timer Interrupt.
D
CCCCCCCCDDDDDD[S] Supervisor Timer Interrupt.
D
CCCCCCCCDDDDDDDD[S] Supervisor Timer Interrupt.
D
CCCCCCCCDDDDDDDD
switch to [PID = 1 PRIORITY = 37 COUNTER = 4]

B
CCCCCCCCDDDDDDDDDB[S] Supervisor Timer Interrupt.
B
CCCCCCCCDDDDDDDDDBB[S] Supervisor Timer Interrupt.
B
CCCCCCCCDDDDDDDDDBBB[S] Supervisor Timer Interrupt.
B
CCCCCCCCDDDDDDDDDBBBB
NR_TASKS = 4, PRIORITY test passed!

```

NR_TASKS = 8

NR_TASKS = 16

```
FFFFFFFFFFFFKKKKKKKKKKGGGGGGGGCCCCCCCCDDDDDDJJJJJJPPPPPEEEEEMMM[S] Supervisor Timer Interrupt.
M
FFFFFFFFFFFFKKKKKKKKKKGGGGGGGGCCCCCCCCDDDDDDJJJJJJPPPPPEEEEEMMM
switch to [PID = 1 PRIORITY = 37 COUNTER = 4]

B
FFFFFFFFFFFFKKKKKKKKKKGGGGGGGGCCCCCCCCDDDDDDJJJJJJPPPPPEEEEEMMMB[S] Supervisor Timer Interrupt.
B
FFFFFFFFFFFFKKKKKKKKKKGGGGGGGGCCCCCCCCDDDDDDJJJJJJPPPPPEEEEEMMMBB[S] Supervisor Timer Interrupt.
B
FFFFFFFFFFFFKKKKKKKKKKGGGGGGGGCCCCCCCCDDDDDDJJJJJJPPPPPEEEEEMMMBBB[S] Supervisor Timer Interrupt.
B
FFFFFFFFFFFFKKKKKKKKKKGGGGGGGGCCCCCCCCDDDDDDJJJJJJPPPPPEEEEEMMMBBBB
switch to [PID = 13 PRIORITY = 40 COUNTER = 3]

N
FFFFFFFFFFFFKKKKKKKKKKGGGGGGGGCCCCCCCCDDDDDDJJJJJJPPPPPEEEEEMMMBBBN[S] Supervisor Timer Interrupt.
N
FFFFFFFFFFFFKKKKKKKKKKGGGGGGGGCCCCCCCCDDDDDDJJJJJJPPPPPEEEEEMMMBBBNN[S] Supervisor Timer Interrupt.
N
FFFFFFFFFFFFKKKKKKKKKKGGGGGGGGCCCCCCCCDDDDDDJJJJJJPPPPPEEEEEMMMBBBNNN
switch to [PID = 11 PRIORITY = 21 COUNTER = 3]

L
FFFFFFFFFFFFKKKKKKKKKKGGGGGGGGCCCCCCCCDDDDDDJJJJJJPPPPPEEEEEMMMBBBNNNL[S] Supervisor Timer Interrupt.
L
FFFFFFFFFFFFKKKKKKKKKKGGGGGGGGCCCCCCCCDDDDDDJJJJJJPPPPPEEEEEMMMBBBNNNLL[S] Supervisor Timer Interrupt.
L
FFFFFFFFFFFFKKKKKKKKKKGGGGGGGGCCCCCCCCDDDDDDJJJJJJPPPPPEEEEEMMMBBBNNNLLL
switch to [PID = 14 PRIORITY = 6 COUNTER = 2]

O
FFFFFFFFFFFFKKKKKKKKKKGGGGGGGGCCCCCCCCDDDDDDJJJJJJPPPPPEEEEEMMMBBBNNNLLLO[S] Supervisor Timer Interrupt.
O
FFFFFFFFFFFFKKKKKKKKKKGGGGGGGGCCCCCCCCDDDDDDJJJJJJPPPPPEEEEEMMMBBBNNNLLLOO
switch to [PID = 7 PRIORITY = 5 COUNTER = 2]

H
FFFFFFFFFFFFKKKKKKKKKKGGGGGGGGCCCCCCCCDDDDDDJJJJJJPPPPPEEEEEMMMBBBNNNLLLOOH[S] Supervisor Timer Interrupt.
H
FFFFFFFFFFFFKKKKKKKKKKGGGGGGGGCCCCCCCCDDDDDDJJJJJJPPPPPEEEEEMMMBBBNNNLLLOOH
switch to [PID = 8 PRIORITY = 25 COUNTER = 1]

I
FFFFFFFFFFFFKKKKKKKKKKGGGGGGGGCCCCCCCCDDDDDDJJJJJJPPPPPEEEEEMMMBBBNNNLLLOOHI
NR_TASKS = 16, PRIORITY test passed!
```

运行 短作业优先调度算法

OpenSBI v0.9



Platform Name

: riscv-virtio,qemu

Platform Features

: timer,mfdeleg

Platform HART Count

: 1

Firmware Base

: 0x80000000

Firmware Size

: 100 KB

Runtime SBI Version

: 0.2

Domain0 Name

: root

Domain0 Boot HART

: 0

Domain0 HARTs

: 0*

```

Domain0 Region00      : 0x0000000080000000-0x000000008001ffff ()
Domain0 Region01      : 0x0000000000000000-0xffffffffffff (R,W,X)
Domain0 Next Address   : 0x0000000080200000
Domain0 Next Arg1      : 0x0000000087000000
Domain0 Next Mode      : S-mode
Domain0 SysReset       : yes

```

```

Boot HART ID          : 0
Boot HART Domain      : root
Boot HART ISA          : rv64imafdcsu
Boot HART Features     : scounteren,mcounteren,time
Boot HART PMP Count    : 16
Boot HART PMP Granularity : 4
Boot HART PMP Address Bits: 54
Boot HART MHPM Count   : 0
Boot HART MHPM Count   : 0
Boot HART MIDELEG      : 0x0000000000000222
Boot HART MEDELEG      : 0x000000000000b109

```

...mm_init done!

...proc_init done!

2023 Hello RISC-V

switch to [PID = 1 COUNTER = 4]

```

[PID = 1] is running. auto_inc_local_var = 1
[COUNTER = 4]
[S] Supervisor Timer Interrupt.
[PID = 1] is running. auto_inc_local_var = 2
[COUNTER = 3]
[S] Supervisor Timer Interrupt.
[PID = 1] is running. auto_inc_local_var = 3
[COUNTER = 2]
[S] Supervisor Timer Interrupt.
[PID = 1] is running. auto_inc_local_var = 4
[COUNTER = 0]

```

switch to [PID = 3 COUNTER = 8]

```

[PID = 3] is running. auto_inc_local_var = 1
[COUNTER = 8]
[S] Supervisor Timer Interrupt.
[PID = 3] is running. auto_inc_local_var = 2
[COUNTER = 7]
[S] Supervisor Timer Interrupt.
[PID = 3] is running. auto_inc_local_var = 3
[COUNTER = 6]
[S] Supervisor Timer Interrupt.
[PID = 3] is running. auto_inc_local_var = 4
[COUNTER = 5]
[S] Supervisor Timer Interrupt.
[PID = 3] is running. auto_inc_local_var = 5
[COUNTER = 4]
[S] Supervisor Timer Interrupt.
[PID = 3] is running. auto_inc_local_var = 6

```

```

[COUNTER = 3]
[S] Supervisor Timer Interrupt.
[PID = 3] is running. auto_inc_local_var = 7
[COUNTER = 2]
[S] Supervisor Timer Interrupt.
[PID = 3] is running. auto_inc_local_var = 8
[COUNTER = 0]

switch to [PID = 2 COUNTER = 9]

[PID = 2] is running. auto_inc_local_var = 1
[COUNTER = 9]
[S] Supervisor Timer Interrupt.
[PID = 2] is running. auto_inc_local_var = 2
[COUNTER = 8]
[S] Supervisor Timer Interrupt.
[PID = 2] is running. auto_inc_local_var = 3
[COUNTER = 7]
[S] Supervisor Timer Interrupt.
[PID = 2] is running. auto_inc_local_var = 4
[COUNTER = 6]
[S] Supervisor Timer Interrupt.
[PID = 2] is running. auto_inc_local_var = 5
[COUNTER = 5]
[S] Supervisor Timer Interrupt.
[PID = 2] is running. auto_inc_local_var = 6
[COUNTER = 4]
[S] Supervisor Timer Interrupt.
[PID = 2] is running. auto_inc_local_var = 7
[COUNTER = 3]
[S] Supervisor Timer Interrupt.
[PID = 2] is running. auto_inc_local_var = 8
[COUNTER = 2]
[S] Supervisor Timer Interrupt.
[PID = 2] is running. auto_inc_local_var = 9
[COUNTER = 0]

switch to [PID = 0 COUNTER = 0]

[S] Supervisor Timer Interrupt.
[S] Supervisor Timer Interrupt.
.....

```

运行 优先级调度算法

Launch the qemu

OpenSBI v0.9

```

/  _  \      /  _  \  _  \  _  \
| |  | | _  _  | ( _  | | ) || |

```



```

| | | | ' _ \ / _ \ ' _ \ \ _ \ | _ < | |
| | | | | ) | _ / | | | _ ) | | ) | | | _
\ _ / | . _ / \ _ | | | _ / | _ / | _ |
| |
|_|

```

```

Platform Name           : riscv-virtio,qemu
Platform Features       : timer,mfdeleg
Platform HART Count     : 1
Firmware Base           : 0x80000000
Firmware Size           : 100 KB
Runtime SBI Version     : 0.2

Domain0 Name            : root
Domain0 Boot HART       : 0
Domain0 HARTs           : 0*
Domain0 Region00        : 0x0000000080000000-0x000000008001ffff ()
Domain0 Region01        : 0x0000000000000000-0xffffffffffffffff (R,W,X)
Domain0 Next Address    : 0x0000000080200000
Domain0 Next Arg1       : 0x0000000087000000
Domain0 Next Mode       : S-mode
Domain0 SysReset        : yes

Boot HART ID            : 0
Boot HART Domain        : root
Boot HART ISA           : rv64imafdcsu
Boot HART Features      : scounteren,mcounteren,time
Boot HART PMP Count     : 16
Boot HART PMP Granularity : 4
Boot HART PMP Address Bits: 54
Boot HART MHPM Count    : 0
Boot HART MHPM Count    : 0
Boot HART MIDELEG       : 0x0000000000000222
Boot HART MEDELEG       : 0x000000000000b109
...mm_init done!
...proc_init done!
2023 Hello RISC-V

```

```
switch to [PID = 2 PRIORITY = 88 COUNTER = 9]
```

```

[PID = 2] is running. auto_inc_local_var = 1
[COUNTER = 9]
[S] Supervisor Timer Interrupt.
[PID = 2] is running. auto_inc_local_var = 2
[COUNTER = 8]
[S] Supervisor Timer Interrupt.
[PID = 2] is running. auto_inc_local_var = 3
[COUNTER = 7]
[S] Supervisor Timer Interrupt.
[PID = 2] is running. auto_inc_local_var = 4
[COUNTER = 6]
[S] Supervisor Timer Interrupt.
[PID = 2] is running. auto_inc_local_var = 5
[COUNTER = 5]

```

```
[S] Supervisor Timer Interrupt.  
[PID = 2] is running. auto_inc_local_var = 6  
[COUNTER = 4]  
[S] Supervisor Timer Interrupt.  
[PID = 2] is running. auto_inc_local_var = 7  
[COUNTER = 3]  
[S] Supervisor Timer Interrupt.  
[PID = 2] is running. auto_inc_local_var = 8  
[COUNTER = 2]  
[S] Supervisor Timer Interrupt.  
  
.....
```

思考题

1.在 RV64 中一共用 32 个通用寄存器，为什么 context_switch 中只保存了14个？

- 首先出于性能和资源管理的考虑。在进行上下文切换时，保存和恢复所有32个通用寄存器的状态可能会导致不必要的开销。
- 在上下文切换过程中，只需要保存和恢复那些在函数调用期间被破坏的寄存器，以确保被调用函数的状态不受影响。这通常包括保存调用者保存寄存器和一些其他需要的寄存器。
- RV64架构通常定义了哪些寄存器是调用者保存的，哪些是被调用者保存的。调用者保存寄存器通常在函数调用前被保存，函数调用结束后再恢复，以保护调用者的上下文。被调用者保存寄存器则由被调用的函数负责保存和恢复。
- 所以，context_switch可能只保存了需要在上下文切换期间保护的寄存器，而其他的寄存器可以在被调用的函数中负责保存和恢复。

2.当线程第一次调用时，其 ra 所代表的返回点是 __dummy。那么在之后的线程调用中 context_switch 中，ra 保存/恢复的函数返回点是什么呢？请同学用 gdb 尝试追踪一次完整的线程切换流程，并关注每一次 ra 的变换(需要截图)。

首先会先进入 trap_handler 函数,此时 ra 记录调用 trap_handler 的下一条指令，然后在 trap_handler 中调用 do_timer 函数，ra 是调用 do_timer 的下一条指令。然后在 do_timer 中因为 current 是 idle，所以要进行调度，调用 schedule 函数，ra 变成调用 schedule 的下一条指令，然后进入 dummy 函数，ra 为 __dummy 的地址，完成第一次调用，然后 ra 逐个恢复，使程序依次回到schedule(),do_timer(),trap_handler(),然后再次调用 trap_handler()。

```
Breakpoint 2, trap_handler (scause=9223372036854775813, sepc=2149584180)
  at trap.c:14
14      unsigned long temp = 1;
(gdb) i r ra
ra      0x80200100      0x80200100 <_traps+152>
(gdb) c
Continuing.

Breakpoint 1, do_timer () at proc.c:98
98      if(current == idle){
(gdb) i r ra
ra      0x80200cc4      0x80200cc4 <trap_handler+80>
(gdb) c
Continuing.

Breakpoint 3, schedule () at proc.c:113
113     uint64 judge = 1;
(gdb) i r ra
ra      0x8020096c      0x8020096c <do_timer+48>
(gdb) c
Continuing.

Breakpoint 4, dummy () at proc.c:60
60      schedule_test();
(gdb) i r ra
ra      0x80200190      0x80200190 <__dummy>
(gdb) c
Continuing.

Breakpoint 2, trap_handler (scause=9223372036854775813, sepc=2149582788)
  at trap.c:14
14      unsigned long temp = 1;
(gdb) i r ra
ra      0x80200100      0x80200100 <_traps+152>
(gdb) c
Continuing.

Breakpoint 1, do_timer () at proc.c:98
98      if(current == idle){
(gdb) i r ra
ra      0x80200cc4      0x80200cc4 <trap_handler+80>
```

```
Breakpoint 2, trap_handler (scause=9223372036854775813, sepc=2149582776)
  at trap.c:14
14      unsigned long temp = 1;
(gdb) i r ra
ra      0x80200100      0x80200100 <_traps+152>
(gdb) c
Continuing.

Breakpoint 1, do_timer () at proc.c:98
98      if(current == idle){
(gdb) i r ra
ra      0x80200cc4      0x80200cc4 <trap_handler+80>
(gdb) c
Continuing.

Breakpoint 2, trap_handler (scause=9223372036854775813, sepc=2149582796)
  at trap.c:14
14      unsigned long temp = 1;
(gdb) i r ra
ra      0x80200100      0x80200100 <_traps+152>
(gdb) c
Continuing.

Breakpoint 1, do_timer () at proc.c:98
98      if(current == idle){
(gdb) i r ra
ra      0x80200cc4      0x80200cc4 <trap_handler+80>
(gdb) c
Continuing.

Breakpoint 2, trap_handler (scause=9223372036854775813, sepc=2149582788)
  at trap.c:14
14      unsigned long temp = 1;
(gdb) i r ra
ra      0x80200100      0x80200100 <_traps+152>
(gdb) c
Continuing.

Breakpoint 1, do_timer () at proc.c:98
98      if(current == idle){
(gdb) i r ra
ra      0x80200cc4      0x80200cc4 <trap_handler+80>
```

第一个进程调用第二个进程过程基本相同，唯一不同的地方在于 do_timer() 中调用 schedule() 的地方不同了，所以 ra 保存略有偏差。

```
(gdb) c
Continuing.

Breakpoint 3, schedule () at proc.c:113
113      uint64 judge = 1;
(gdb) i r ra
ra      0x80200994      0x80200994 <do_timer+88>
(gdb) c
Continuing.

Breakpoint 4, dummy () at proc.c:60
60      schedule_test();
(gdb) i r ra
ra      0x80200190      0x80200190 <__dummy>
(gdb) c
Continuing.

Breakpoint 2, trap_handler (scause=9223372036854775813, sepc=2149582800)
at trap.c:14
14      unsigned long temp = 1;
(gdb) i r ra
ra      0x80200100      0x80200100 <_traps+152>
(gdb) █
```