

《操作系统原理》实验报告

姓名：韩艺轩
学院：计算机科学与技术学院
专业：计算机科学与技术（图灵班）
邮箱：2674361965@qq.com
指导教师：申文博
报告日期：2023/11/20

- Lab3-RV64 虚拟内存管理
 - 实验步骤
 - 准备工作
 - 开启虚拟内存映射
 - 实现 `setup_vm`
 - 实现 `setup_vm_final`
 - 实现 `create_mapping`
 - 实验结果
 - 思考题

Lab3-RV64 虚拟内存管理

实验步骤

准备工作

- 在 defs.h 添加如下内容：

```
#define OPENSBI_SIZE (0x200000)

#define VM_START (0xffffffe000000000)
#define VM_END   (0xfffffffff0000000)
#define VM_SIZE   (VM_END - VM_START)

#define PA2VA_OFFSET (VM_START - PHY_START)
```

- 从 repo 同步以下代码: vmlinux.lds。并按照以下步骤将这些文件正确放置。

```
.
├── arch
│   └── riscv
│       └── kernel
│           └── vmlinux.lds
```

新的链接脚本中的 ramv 代表 VMA (Virtual Memory Address) 即虚拟地址, ram 则代表 LMA (Load Memory Address), 即我们 OS image 被 load 的地址, 可以理解为物理地址。使用以上的 vmlinux.lds 进行编译之后, 得到的 System.map 以及 vmlinux 中的符号采用的都是虚拟地址, 方便之后 Debug。

- 从本实验开始我们需要使用刷新缓存的指令扩展, 并自动在编译项目前执行 clean 任务来防止对头文件的修改无法触发编译任务。在项目顶层目录的 Makefile 中需要做如下更改:

```
# Makefile
...
ISA=rv64imafd_zifencei
...
all: clean
    ${MAKE} -C lib all
    ${MAKE} -C test all
    ${MAKE} -C init all
    ${MAKE} -C arch/riscv all
    @echo -e '\n'Build Finished OK
...
```

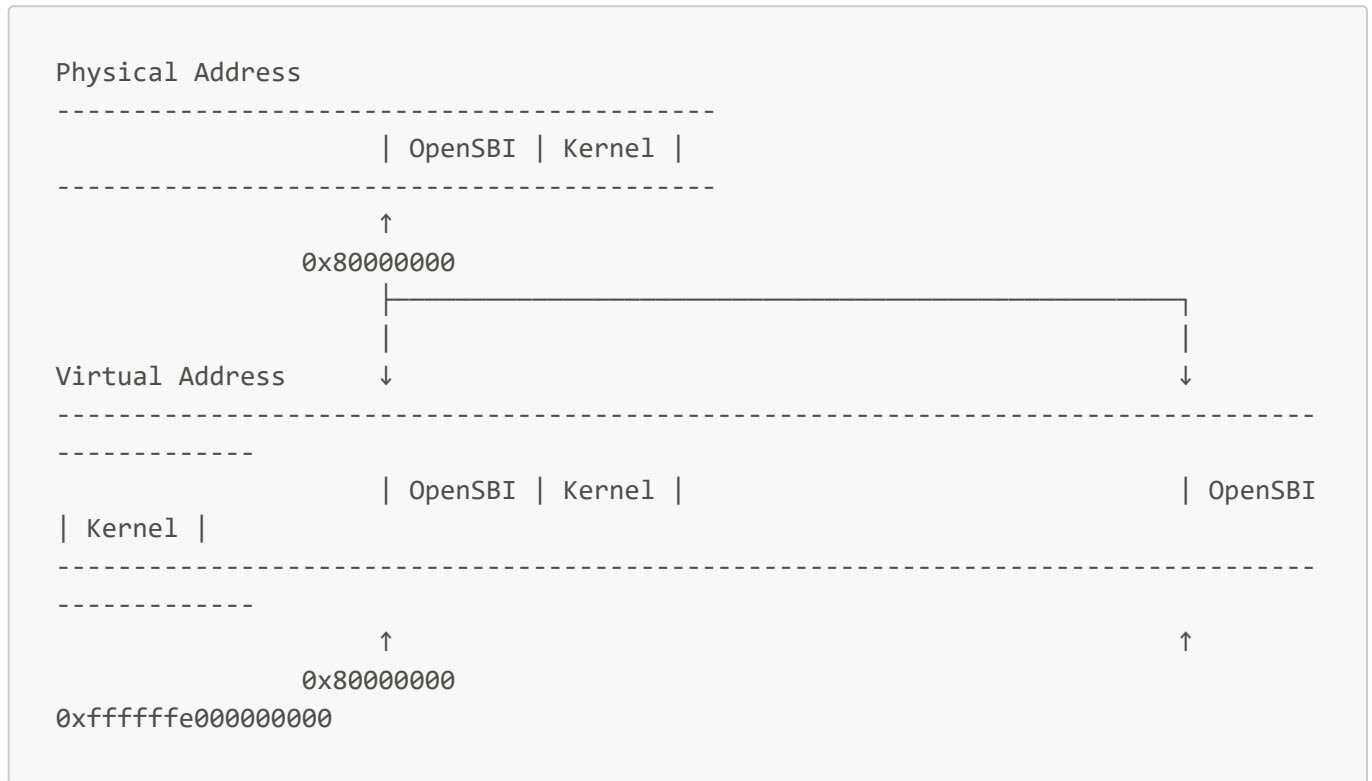
开启虚拟内存映射

在 RISC-V 中开启虚拟地址被分为了两步：setup_vm 以及 setup_vm_final，下面将介绍相关的具体实现。

实现 setup_vm

将 `0x80000000` 开始的 1GB 区域进行两次映射，其中一次是等值映射 ($PA == VA$)，另一次是将其映射到 `direct mapping area` (使得 $PA + PV2VA_OFFSET == VA$)。

对于函数 `setup_vm` 的实现，我们需要根据如下的两次映射将 `0x80200000` 开始的 1GB 区域进行两次映射：



因为要映射 1GB 的区域，因为 $1GB = 2^{30}bit$ ，所以 `vpn2` 右边的 30 位刚好可以表示 1GB 的空间，所以可以将 `vpn2` 作为 index，将物理地址改写成 entry 的格式填入顶级页表中，物理地址与 `0x0000007fc0000000` 做与操作就可以使除 `vpn2` 的其他位变为 0。Page Table Entry 的权限 `V | R | W | X` 位设置为 1，与 0xf 做或操作可以设置权限。

```
// arch/riscv/kernel/vm.c
unsigned long early_pgtbl[512] __attribute__((__aligned__(0x1000)));

void setup_vm() {
    // printk("Setup_vm...\n");
    memset(early_pgtbl, 0x0, PGSIZE);
    uint64 vpn2 = (PHY_START & 0x0000007fc0000000) >> 30;
    early_pgtbl[vpn2] = ((PHY_START >> 30) << 28) | 0xf;
    vpn2 = (VM_START & 0x0000007fc0000000) >> 30;
    early_pgtbl[vpn2] = ((PHY_START >> 30) << 28) | 0xf;
    // printk("leaving setup_vm...\n");
    return ;
}
```

完成上述映射之后，通过 `relocate` 函数，完成对 `satp` 的设置，以及跳转到对应的虚拟地址。
`0xffffffffdf80000000` 是 `VM_START - PHY_START` 的值，在 `relocate` 中，我们将 `ra` 和 `sp` 都映射到虚拟地址。
 对于 `early_pgtbl`，它的地址目前是虚拟地址，但是要将物理地址写入 `satp`，所以需要减去
`0xffffffffdf80000000`

```
// head.S

.global relocate
relocate:

    li t0, 0xffffffffdf80000000
    add ra, ra, t0
    add sp, sp, t0

    la t0, early_pgtbl
    li t1, 0xffffffffdf80000000
    sub t0, t0, t1
    srli t0, t0, 12
    li t1, 8
    slli t1, t1, 60
    or t0, t0, t1
    csrw satp, t0

    sfence.vma zero, zero

    fence.i

    ret
```

至此我们已经完成了虚拟地址的开启，之后我们运行的代码也都将在虚拟地址上运行。

实现 `setup_vm_final`

由于 `setup_vm_final` 中需要申请页面的接口，应该在其之前完成内存管理初始化，需要修改 `mm.c` 中的代码，`mm.c` 中初始化的函数接收的起始结束地址需要调整为虚拟地址。

```
void mm_init(void) {
    kfreerange(_kernel, (char *) (PHY_END + PA2VA_OFFSET));
    printk("...mm_init done!\n");
}
```

现在代码已经在虚拟地址运行，完成 `setup_vm_final` 将完成对 `kernel` 的映射。函数内主要是对 `create_mapping` 的调用，完善三级页表的构建。然后修改 `satp`，同样要修改成物理地址后再存入 `satp`。

```
// arch/riscv/kernel/vm.c
extern uint64 _stext;
extern uint64 _etext;
```

```

extern uint64 _srodata;
extern uint64 _erodata;
extern uint64 _sdata;
extern uint64 _edata;
unsigned long swapper_pg_dir[512] __attribute__((__aligned__(0x1000)));

void setup_vm_final() {
    // printk("Setup_vm_final...\n");
    memset(swapper_pg_dir, 0x0, PGSIZE);
    // No OpenSBI mapping required

    // mapping kernel text X|-|R|V
    uint64 va = VM_START + OPENSBI_SIZE;
    uint64 pa = PHY_START + OPENSBI_SIZE;
    create_mapping(swapper_pg_dir, va, pa, (uint64)_srodata-(uint64)_stext, 0xb);

    // mapping kernel rodata -|-|R|V
    va += (uint64)_srodata-(uint64)_stext;
    pa += (uint64)_srodata-(uint64)_stext;
    create_mapping(swapper_pg_dir, va, pa, (uint64)_sdata-(uint64)_srodata, 0x3);

    // mapping other memory -|W|R|V
    va += (uint64)_sdata-(uint64)_srodata;
    pa += (uint64)_sdata-(uint64)_srodata;
    create_mapping(swapper_pg_dir, va, pa, PHY_SIZE - ((uint64)_sdata -
(uint64)_stext), 0x7);

    // set satp with swapper_pg_dir
    uint64 table = (((uint64)swapper_pg_dir - PA2VA_OFFSET) >> 12) | (8UL << 60);
    // printk("before asm...\n");
    __asm__ volatile (
        "csrw satp, %[table]\n"
        :
        : [table] "r" (table)
        : "memory"
    );

    // flush TLB
    asm volatile("sfence.vma zero, zero\n");

    // flush icache
    asm volatile("fence.i\n");
    printk("leaving setup_vm_final...\n");
    return ;
}

```

实现create_mapping

本函数主要的目的是构建三级页表，写入正确的page entry来完成正确的映射。物理地址和虚拟地址的offset是12位，刚好对应每页4KB的内存，也就是说第三级页表的每一个entry都对应一页，所以我们需要将给定size的内存一页一页地完成映射。对于新建页表，使用kalloc()新建一页来作为页表，但新建的页表的地址现在是虚拟地址，需要转化成物理地址才能作为PTE写入上一级页表。

```
// arch/riscv/kernel/vm.c

void create_mapping(uint64 *pgtbl, uint64 va, uint64 pa, uint64 sz, uint64 perm) {
    // printk("create_mapping...\n");
    uint64 *second;
    uint64 *third;

    for(uint64 i = va; i < va + sz; i = i + PGSIZE){
        uint64 vpn2 = (i >> 30) & 0x1ff;
        uint64 vpn1 = (i >> 21) & 0x1ff;
        uint64 vpn0 = (i >> 12) & 0x1ff;
        if(!(pgtbl[vpn2] & 0x1)){
            second = (uint64 *)kalloc();
            pgtbl[vpn2] = (((uint64)second - PA2VA_OFFSET) >> 12) << 10 | 0x1;
            second = (uint64 *)((pgtbl[vpn2] >> 10) << 12);
            memset(second, 0x0, PGSIZE);
        }else{
            second = (uint64 *)((pgtbl[vpn2] >> 10) << 12);
        }
        if(!(second[vpn1] & 0x1)){
            third = (uint64 *)kalloc();
            second[vpn1] = (((uint64)third - PA2VA_OFFSET) >> 12) << 10 | 0x1;
            third = (uint64 *)((second[vpn1] >> 10) << 12);
            memset(third, 0x0, PGSIZE);
        }else{
            third = (uint64 *)((second[vpn1] >> 10) << 12);
        }
        third[vpn0] = ((pa >> 12) << 10) | perm;
        pa += PGSIZE;
    }

    return ;
}
```

实验结果

编译后内核可以想lab2一样运行。

```
Boot HART MHPM Count      : 0
Boot HART MIDELEG         : 0x0000000000000222
Boot HART MEDELEG         : 0x000000000000b109
...mm_init done!
leaving setup_vm_final...
...proc_init done!
2023 Hello RISC-V

switch to [PID = 8 COUNTER = 1]

[PID = 8] is running. auto_inc_local_var = 1
[COUNTER = 0]

switch to [PID = 7 COUNTER = 2]

[PID = 7] is running. auto_inc_local_var = 1
[COUNTER = 2]
[S] Supervisor Timer Interrupt.
[PID = 7] is running. auto_inc_local_var = 2
[COUNTER = 0]

switch to [PID = 14 COUNTER = 2]

[PID = 14] is running. auto_inc_local_var = 1
[COUNTER = 2]
[S] Supervisor Timer Interrupt.
[PID = 14] is running. auto_inc_local_var = 2
[COUNTER = 0]

switch to [PID = 11 COUNTER = 3]

[PID = 11] is running. auto_inc_local_var = 1
[COUNTER = 3]
[S] Supervisor Timer Interrupt.
[PID = 11] is running. auto_inc_local_var = 2
[COUNTER = 2]
[S] Supervisor Timer Interrupt.
[PID = 11] is running. auto_inc_local_var = 3
[COUNTER = 0]

switch to [PID = 13 COUNTER = 3]

[PID = 13] is running. auto_inc_local_var = 1
[COUNTER = 3]
[S] Supervisor Timer Interrupt.
[PID = 13] is running. auto_inc_local_var = 2
[COUNTER = 2]
[S] Supervisor Timer Interrupt.
```

思考题

1. 验证 `.text`, `.rodata` 段的属性是否成功设置, 给出截图。

Linux内核可以正常运行, 代表`.text`段的可执行的属性就已经得到验证。验证`text`段不可写的属性。在函数`start_kernel`中修改`text`的值, 在调用`gdb`的时候发现在这一行卡住了, 无法运行, 验证不可写。


```

Domain0 Name      : root
Domain0 Boot HART : 0
Domain0 HARTs     : 0*
Domain0 Region00  : 0x0000000080000000-0x000000008001ffff ( )
Domain0 Region01  : 0x0000000000000000-0xffffffffffffff (R,W,X)
Domain0 Next Address : 0x0000000080200000
Domain0 Next Arg1  : 0x0000000087000000
Domain0 Next Mode  : S-mode
Domain0 SysReset   : yes

Boot HART ID      : 0
Boot HART Domain  : root
Boot HART ISA     : rv64imafdcu
Boot HART Features : scounteren,mcounteren,time
Boot HART PMP Count : 16
Boot HART PMP Granularity : 4
Boot HART PMP Address Bits: 54
Boot HART MHPM Count : 0
Boot HART MHPM Count : 0
Boot HART MIDELEG : 0x0000000000000222
Boot HART MEDELEG : 0x000000000000b109
...mm_init done!
leaving setup_vm_final...
...proc_init done!
2023 Hello RISC-V

switch to [PID = 8 COUNTER = 1]

[PID = 8] is running. auto_inc_local_var = 1
[COUNTER = 0]

switch to [PID = 7 COUNTER = 2]

[PID = 7] is running. auto_inc_local_var = 1
[COUNTER = 2]
[S] Supervisor Timer Interrupt.
[PID = 7] is running. auto_inc_local_var = 2
[COUNTER = 0]

switch to [PID = 14 COUNTER = 2]

[PID = 14] is running. auto_inc_local_var = 1
[COUNTER = 2]
[S] Supervisor Timer Interrupt.
[PID = 14] is running. auto_inc_local_var = 2
[COUNTER = 0]

(gdb) quit
hanyixuan@DESKTOP-LP7R135:~/os23fall-stu/src/lab3$ gdb-multiarch ./vmlin
ux
GNU gdb (Ubuntu 12.1-0ubuntu1~22.04) 12.1
Copyright (C) 2022 Free Software Foundation, Inc.
License GPLv3+: GNU GPL version 3 or later <http://gnu.org/licenses/gpl.
html>
This is free software: you are free to change and redistribute it.
There is NO WARRANTY, to the extent permitted by law.
Type "show copying" and "show warranty" for details.
This GDB was configured as "x86_64-linux-gnu".
Type "show configuration" for configuration details.
For bug reporting instructions, please see:
<https://www.gnu.org/software/gdb/bugs/>.
Find the GDB manual and other documentation resources online at:
<http://www.gnu.org/software/gdb/documentation/>.

For help, type "help".
Type "apropos word" to search for commands related to "word"...
Reading symbols from ./vmlinux...
(gdb) target remote :1234
Remote debugging using :1234
0x0000000000001000 in ?? ()
(gdb) b start_kernel
Breakpoint 1 at 0xffffffe0002011b0: file main.c, line 9.
(gdb) c
Continuing.

Breakpoint 1, start_kernel () at main.c:9
9     printk("2023");
(gdb) list
4     extern void test();
5     extern char _stext[];
6
7     int start_kernel()
8     {
9         printk("2023");
10        printk(" Hello RISC-V\n");
11
12        _stext[0] = (uint64)0x1;
13        printk("%lu\n", _stext[0]);
(gdb) n
10        printk(" Hello RISC-V\n");
(gdb) n
12        _stext[0] = (uint64)0x1;
(gdb) n

```

验证rodata只读的属性。在函数start_kernel中修改rodata的值，在调用gdb的时候发现在这一行卡住了，无法运行，验证不可写。

```

Firmware Base      : 0x80000000
Firmware Size      : 100 KB
Runtime SBI Version : 0.2

Domain0 Name      : root
Domain0 Boot HART : 0
Domain0 HARTs     : 0*
Domain0 Region00  : 0x0000000080000000-0x000000008001ffff ( )
Domain0 Region01  : 0x0000000000000000-0xffffffffffffff (R,W,X)
Domain0 Next Address : 0x0000000080200000
Domain0 Next Arg1  : 0x0000000087000000
Domain0 Next Mode  : S-mode
Domain0 SysReset   : yes

Boot HART ID      : 0
Boot HART Domain  : root
Boot HART ISA     : rv64imafdcu
Boot HART Features : scounteren,mcounteren,time
Boot HART PMP Count : 16
Boot HART PMP Granularity : 4
Boot HART PMP Address Bits: 54
Boot HART MHPM Count : 0
Boot HART MHPM Count : 0
Boot HART MIDELEG : 0x0000000000000222
Boot HART MEDELEG : 0x000000000000b109
...mm_init done!
leaving setup_vm_final...
...proc_init done!
2023 Hello RISC-V

switch to [PID = 8 COUNTER = 1]

[PID = 8] is running. auto_inc_local_var = 1
[COUNTER = 0]

switch to [PID = 7 COUNTER = 2]

[PID = 7] is running. auto_inc_local_var = 1
[COUNTER = 2]
[S] Supervisor Timer Interrupt.
[PID = 7] is running. auto_inc_local_var = 2
[COUNTER = 0]

switch to [PID = 14 COUNTER = 2]

[PID = 14] is running. auto_inc_local_var = 1
[COUNTER = 2]
[

hanyixuan@DESKTOP-LP7R135:~/os23fall-stu/src/lab3$ gdb-multiarch ./vmlin
ux
GNU gdb (Ubuntu 12.1-0ubuntu1~22.04) 12.1
Copyright (C) 2022 Free Software Foundation, Inc.
License GPLv3+: GNU GPL version 3 or later <http://gnu.org/licenses/gpl.
html>
This is free software: you are free to change and redistribute it.
There is NO WARRANTY, to the extent permitted by law.
Type "show copying" and "show warranty" for details.
This GDB was configured as "x86_64-linux-gnu".
Type "show configuration" for configuration details.
For bug reporting instructions, please see:
<https://www.gnu.org/software/gdb/bugs/>.
Find the GDB manual and other documentation resources online at:
<http://www.gnu.org/software/gdb/documentation/>.

For help, type "help".
Type "apropos word" to search for commands related to "word"...
Reading symbols from ./vmlinux...
(gdb) target remote :1234
Remote debugging using :1234
0x0000000000001000 in ?? ()
(gdb) b start_kernel
Breakpoint 1 at 0xffffffe0002011b0: file main.c, line 9.
(gdb) c
Continuing.

Breakpoint 1, start_kernel () at main.c:9
9     printk("2023");
(gdb) list
4     extern void test();
5     extern char _srodata[];
6
7     int start_kernel()
8     {
9         printk("2023");
10        printk(" Hello RISC-V\n");
11
12        _srodata[0] = (uint64)0x1;
13        printk("%lu\n", _srodata[0]);
(gdb) n
10        printk(" Hello RISC-V\n");
(gdb) n
12        _srodata[0] = (uint64)0x1;
(gdb) n

```

2. 为什么我们在`setup_vm`中需要做等值映射?

因为我们在创建三级页表以及设置`satp`的时候都需要将物理地址作为PTE存储，并且需要通过物理地址访问各级页表，如果我们不做等值映射，就会出现访问错误等情况。

3. 在Linux中，是不需要做等值映射的。请探索一下不在`setup_vm`中做等值映射的方法。

在`setup_vm`中注释掉等值映射的部分，在`relocate`中会在`sfence`这条指令时访存错误。将`relocate`改写成这样，在一开始设置`stvec`在访存错误时会直接跳到虚拟地址的`sfence`处，然后继续向下执行。

```
relocate:
    li t0, 0xfffffffffe002000e8
    csrw stvec, t0

    csrr t0, sie
    ori t0, t0, 1<<5
    csrw sie, t0

    csrr t0, sstatus
    ori t0, t0, 1<<1
    csrw sstatus, t0

    li t0, 0xffffffffdf80000000
    add ra, ra, t0
    add sp, sp, t0

    la t0, early_pgtbl
    li t1, 0xffffffffdf80000000
    sub t0, t0, t1
    srli t0, t0, 12
    li t1, 8
    slli t1, t1, 60
    or t0, t0, t1
    csrw satp, t0

    sfence.vma zero, zero

    fence.i

    ret
```

因为没有了直接映射，所以我们在`create_mapping`中也不能将页表储存在物理地址了。将`second`和`third`都加上了`PA2VA_OFFSET`。

```
void create_mapping(uint64 *pgtbl, uint64 va, uint64 pa, uint64 sz, uint64 perm) {
    // printk("create_mapping...\n");
    uint64 *second;
    uint64 *third;

    for(uint64 i = va; i < va + sz; i = i + PGSIZE){
```

```
uint64 vpn2 = (i >> 30) & 0x1ff;
uint64 vpn1 = (i >> 21) & 0x1ff;
uint64 vpn0 = (i >> 12) & 0x1ff;
if(!(pgtbl[vpn2] & 0x1)){
    second = (uint64 *)kalloc();
    pgtbl[vpn2] = (((uint64)second - PA2VA_OFFSET) >> 12) << 10 | 0x1;
    second = (uint64 *)(((pgtbl[vpn2] >> 10) << 12) + PA2VA_OFFSET);
    memset(second, 0x0, PGSIZE);
}else{
    second = (uint64 *)(((pgtbl[vpn2] >> 10) << 12) + PA2VA_OFFSET);
}
if(!(second[vpn1] & 0x1)){
    third = (uint64 *)kalloc();
    second[vpn1] = (((uint64)third - PA2VA_OFFSET) >> 12) << 10 | 0x1;
    third = (uint64 *)(((second[vpn1] >> 10) << 12) + PA2VA_OFFSET);
    memset(third, 0x0, PGSIZE);
}else{
    third = (uint64 *)(((second[vpn1] >> 10) << 12) + PA2VA_OFFSET);
}
third[vpn0] = ((pa >> 12) << 10) | perm;
pa += PGSIZE;
}

return ;
}
```

这样不直接映射也可以正常运行了。