

# 《操作系统原理》实验报告

姓名：韩艺轩  
学院：计算机科学与技术学院  
专业：计算机科学与技术（图灵班）  
邮箱：2674361965@qq.com  
指导教师：申文博  
报告日期：2023/12/21

- Lab5 - RV64 缺页异常处理
  - 实验步骤
    - 准备工作
    - 实现VMA
      - 实现 Demand Paging
      - 实现 Page Fault 的检测与处理
  - 实验结果
  - 思考题

# Lab5 - RV64 缺页异常处理

---

## 实验步骤

### 准备工作

从 repo 同步以下文件夹: user 并按照以下步骤将这些文件正确放置。

```
.
├── user
│   ├── Makefile
│   ├── getpid.c
│   ├── link.lds
│   ├── printf.c
│   ├── start.S
│   ├── stddef.h
│   ├── stdio.h
│   ├── syscall.h
│   └── uapp.S
```

### 实现VMA

#### 实现 Demand Paging

修改 `proc.h`, 增加如下相关结构:

```
#define VM_X_MASK      0x0000000000000008
#define VM_W_MASK      0x0000000000000004
#define VM_R_MASK      0x0000000000000002
#define VM_ANONYM      0x0000000000000001

struct vm_area_struct {
    uint64_t vm_start;
    uint64_t vm_end;
    uint64_t vm_flags;
    uint64_t vm_content_offset_in_file;
    uint64_t vm_content_size_in_file;
};

struct task_struct {
    uint64_t state;
    uint64_t counter;
    uint64_t priority;
    uint64_t pid;

    struct thread_struct thread;
    pagetable_t pgd;
};
```

```
uint64_t vma_cnt;
struct vm_area_struct vmas[0];
};
```

为了减少 task 初始化时的开销，我们对一个 Segment 或者 用户态的栈 只需分别建立一个 VMA。

首先要实现一个新建VMA的函数---do\_mmap.

```
void do_mmap(struct task_struct *task, uint64 addr, uint64 length, uint64 flags,
uint64 vm_content_offset_in_file, uint64 vm_content_size_in_file){
    struct vm_area_struct temp;
    temp.vm_start = addr;
    temp.vm_end = addr + length;
    temp.vm_flags = flags;
    temp.vm_content_offset_in_file = vm_content_offset_in_file;
    temp.vm_content_size_in_file = vm_content_size_in_file;
    task->vmas[task->vma_cnt++] = temp;
}
```

修改 load\_program 函数代码，更改为 Demand Paging。Lab4 中已经实现了load\_program,新建VMA的操作放在其中。值得注意的是，Segment中的p\_flags是R-W-X，而VMA需要的是X-W-R-A,所以需要调整一下bit的顺序。

```
uint64 p = 0UL;
if(phdr->p_flags & (1UL) == 1) p = p | VM_X_MASK;
if((phdr->p_flags & (1UL << 1)) >> 1 == 1) p = p | VM_W_MASK;
if((phdr->p_flags & (1UL << 2)) >> 2 == 1) p = p | VM_R_MASK;
do_mmap(task, phdr->p_vaddr, phdr->p_memsz, p, phdr->p_offset, phdr->p_filesz);
```

对于用户栈来说，需要设置权限为VM\_ANONYM | VM\_W\_MASK | VM\_R\_MASK

```
do_mmap(task, USER_END-PGSIZE, PGSIZE, VM_ANONYM | VM_W_MASK | VM_R_MASK, 0, 0);
```

对两个区域建立VMA:

- 代码和数据区域：该区域从 ELF 给出的 Segment 起始用户态虚拟地址 phdr->p\_vaddr 开始，对应文件中偏移量为 phdr->p\_offset 开始的部分
- 用户栈：范围为 [USER\_END - PGSIZE, USER\_END)，权限为 VM\_READ | VM\_WRITE, 并且是匿名的区域。

```
static uint64_t load_program(struct task_struct* task) {
    Elf64_Ehdr* ehdr = (Elf64_Ehdr*)_sramdisk;

    uint64_t phdr_start = (uint64_t)ehdr + ehdr->e_phoff;
```

```
int phdr_cnt = ehdr->e_phnum;

Elf64_Phdr* phdr;
int load_phdr_cnt = 0;
for (int i = 0; i < phdr_cnt; i++) {
    phdr = (Elf64_Phdr*)(phdr_start + sizeof(Elf64_Phdr) * i);
    if (phdr->p_type == PT_LOAD) {
        uint64 p = 0UL;
        if(phdr->p_flags & (1UL) == 1) p = p | VM_X_MASK;
        if((phdr->p_flags & (1UL << 1)) >> 1 == 1) p = p | VM_W_MASK;
        if((phdr->p_flags & (1UL << 2)) >> 2 == 1) p = p | VM_R_MASK;
        do_mmap(task, phdr->p_vaddr, phdr->p_memsz, p, phdr->p_offset, phdr->p_filesz);
        load_phdr_cnt++;
    }
}

do_mmap(task, USER_END-PGSIZE, PGSIZE, VM_ANONYM | VM_W_MASK | VM_R_MASK, 0, 0);
task->thread.sepc = ehdr->e_entry;
task->thread.sstatus = csr_read(sstatus);
task->thread.sstatus = task->thread.sstatus & ~(1 << 8);
task->thread.sstatus = task->thread.sstatus | (1 << 5);
task->thread.sstatus = task->thread.sstatus | (1 << 18);
task->thread.sscratch = USER_END;
return load_phdr_cnt;
}
```

实现 Page Fault 的检测与处理

对于Page Fault,有三种不同的Page Fault.对于scause的值来说。

| Interrupt | Exception Code | Description            |
|-----------|----------------|------------------------|
| 0         | 0xc            | Instruction Page Fault |
| 0         | 0xd            | Load Page Fault        |
| 0         | 0xf            | Store/AMO Page Fault   |

- 修改 trap.c, 添加捕获 Page Fault 的逻辑.遇到Page Fault,就转到do\_page\_fault处理它。

```
void trap_handler(unsigned long scause, unsigned long sepc, struct pt_regs *regs)
{
    unsigned long temp = 1;
    if(scause == 0x8000000000000005){
        clock_set_next_event();

        do_timer();

        printk("[S] Supervisor Timer Interrupt.\n");
    }else if(scause == 8UL){
```

```

        syscall(regs);
        printk("[U] User Environment Call.\n");
    }else if(scause == 0xc | scause == 0xd | scause == 0xf){
        printk("[S] Supervisor Page Falut, ");
        printk("scause: %lx, ", scause);
        printk("stval: %lx, ", regs->stval);
        printk("sepc: %lx\n", regs->sepc);
        do_page_fault(regs);
    }else{
        printk("[S] Unhandled trap, ");
        printk("scause: %lx, ", scause);
        printk("stval: %lx, ", regs->stval);
        printk("sepc: %lx\n", regs->sepc);
        while (1);
    }
}

```

然后还需要实现 `find_vma` 来查找某个虚拟地址是否在某个 `vma` 中。

```

struct vm_area_struct *find_vma(struct task_struct *task, uint64 addr){
    uint64 num = task->vma_cnt;
    // printk("%lx\n", num);
    // while(1);
    for(uint64 i = 0; i < num; i++){
        // printk("%lx\n", task->vmas[i].vm_start);
        if(task->vmas[i].vm_start <= addr && task->vmas[i].vm_end >= addr){
            return &(task->vmas[i]);
        }
    }

    return NULL;
}

```

最后，再正确地实现 `do_page_fault` 就可以了。在 `do_page_fault` 中首先通过 `stval` 获得访问出错的虚拟内存地址 (Bad Address)，然后通过 `find_vma()` 查找 Bad Address 是否在某个 `vma` 中，分配一个页，将这个页映射到对应的用户地址空间，再通过 `(vma->vm_flags & VM_ANONYM)` 获得当前的 VMA 是否是匿名空间，根据 VMA 匿名与否决定将新的页清零或是拷贝 `uapp` 中的内容，如果是要向新的页写入内容，还要注意内容的对齐，将 `stval` 所在的那一页对应的文件中的内容写入新的页中，如果有超出 `[vaddr, vaddr + file_size]` 的部分要置零，且注意页内的对齐。

```

void do_page_fault(struct pt_regs *regs) {
    uint64 stvall = regs->stval;
    struct vm_area_struct *vma_temp = find_vma(current, stvall);
    if(vma_temp != NULL){
        uint64 seg_start = ((uint64)_sramdisk + vma_temp-
>vm_content_offset_in_file);
        uint64 PG_start = (stvall / PGSIZE) * PGSIZE;
        uint64 PG_end = PG_start + PGSIZE;
    }
}

```

```

        char *temp = (char *)alloc_pages(1);
        memset((void*)temp, 0x0, PGSIZE);
        // printk("type:0.\n");
        // while (1);
        if(!(vma_temp->vm_flags & VM_ANONYM)){
            if(PG_start <= vma_temp->vm_start){
                uint64 pre_offset = vma_temp->vm_start - PG_start;
                if(PG_end < vma_temp->vm_start + vma_temp-
>vm_content_size_in_file){
                    // printk("type:1.\n");
                    // while (1);
                    memcpy((void *)((uint64)temp + pre_offset), (void*)
(seg_start), PGSIZE - pre_offset);
                }else{
                    uint64 la_offset = PG_end - (vma_temp->vm_start +
vma_temp->vm_content_size_in_file);
                    memcpy((void *)((uint64)temp + pre_offset), (void*)
(seg_start), PGSIZE - pre_offset - la_offset);
                }
            }else{
                uint64 offset = PG_start - vma_temp->vm_start;
                if(PG_end < vma_temp->vm_start + vma_temp-
>vm_content_size_in_file){
                    // printk("type:3.\n");
                    // while (1);
                    memcpy((void*)(temp), (void*)(seg_start + offset),
PGSIZE);
                }else{
                    // printk("type:4.\n");
                    // while (1);
                    uint64 la_offset = PG_end - (vma_temp->vm_start +
vma_temp->vm_content_size_in_file);
                    if(PG_start<vma_temp->vm_start + vma_temp-
>vm_content_size_in_file)
                        memcpy((void*)(temp), (void*)(seg_start + offset),
PGSIZE - la_offset);
                }
            }
        }
        printk("pgd is %lx\n", vma_temp->vm_flags | 1UL | 1UL << 4);
        create_mapping(current->pgd, (uint64)PG_start, (uint64)temp-
PA2VA_OFFSET, PGSIZE, vma_temp->vm_flags | 1UL | 1UL << 4);
    }else{
        // printk("type:-1.\n");
        // while (1);
    }
}

```

## 实验结果

下面的截图中每个进程被调用了两次，main作为uapp的时候，每个进程会发生**3次**Page Fault.

2023 Hello RISC-V

switch to [PID = 1 COUNTER = 4]

```
[S] Supervisor Page Falut, scause: 00000000000000c, stval: 0000000000100e8, sepc: 0000000000100e8
pgd is 000000000000001f
[S] Supervisor Page Falut, scause: 00000000000000f, stval: 0000003fffffffff8, sepc: 000000000010124
pgd is 0000000000000017
[S] Supervisor Page Falut, scause: 00000000000000d, stval: 000000000011880, sepc: 000000000010140
pgd is 000000000000001f
[PID = 1] is running, variable: 0
[PID = 1] is running, variable: 1
[PID = 1] is running, variable: 2
[S] Supervisor Timer Interrupt.
[PID = 1] is running, variable: 3
[PID = 1] is running, variable: 4
[S] Supervisor Timer Interrupt.
[PID = 1] is running, variable: 5
[PID = 1] is running, variable: 6
[S] Supervisor Timer Interrupt.
[PID = 1] is running, variable: 7
[PID = 1] is running, variable: 8
```

switch to [PID = 4 COUNTER = 5]

```
[S] Supervisor Page Falut, scause: 00000000000000c, stval: 0000000000100e8, sepc: 0000000000100e8
pgd is 000000000000001f
[S] Supervisor Page Falut, scause: 00000000000000f, stval: 0000003fffffffff8, sepc: 000000000010124
pgd is 0000000000000017
[S] Supervisor Page Falut, scause: 00000000000000d, stval: 000000000011880, sepc: 000000000010140
pgd is 000000000000001f
[PID = 4] is running, variable: 0
[PID = 4] is running, variable: 1
[PID = 4] is running, variable: 2
[S] Supervisor Timer Interrupt.
[PID = 4] is running, variable: 3
[PID = 4] is running, variable: 4
[S] Supervisor Timer Interrupt.
[PID = 4] is running, variable: 5
[PID = 4] is running, variable: 6
[S] Supervisor Timer Interrupt.
[PID = 4] is running, variable: 7
[PID = 4] is running, variable: 8
[S] Supervisor Timer Interrupt.
[PID = 4] is running, variable: 9
[PID = 4] is running, variable: 10
```

switch to [PID = 3 COUNTER = 8]

switch to [PID = 3 COUNTER = 8]

```
[S] Supervisor Page Falut, scause: 00000000000000c, stval: 0000000000100e8, sepc: 0000000000100e8
pgd is 000000000000001f
[S] Supervisor Page Falut, scause: 00000000000000f, stval: 0000003fffffffff8, sepc: 000000000010124
pgd is 0000000000000017
[S] Supervisor Page Falut, scause: 00000000000000d, stval: 000000000011880, sepc: 000000000010140
pgd is 000000000000001f
[PID = 3] is running, variable: 0
[PID = 3] is running, variable: 1
[PID = 3] is running, variable: 2
[S] Supervisor Timer Interrupt.
[PID = 3] is running, variable: 3
[PID = 3] is running, variable: 4
[S] Supervisor Timer Interrupt.
[PID = 3] is running, variable: 5
[PID = 3] is running, variable: 6
[S] Supervisor Timer Interrupt.
[PID = 3] is running, variable: 7
[PID = 3] is running, variable: 8
[S] Supervisor Timer Interrupt.
[PID = 3] is running, variable: 9
[PID = 3] is running, variable: 10
[S] Supervisor Timer Interrupt.
[PID = 3] is running, variable: 11
[PID = 3] is running, variable: 12
[S] Supervisor Timer Interrupt.
[PID = 3] is running, variable: 13
[PID = 3] is running, variable: 14
[S] Supervisor Timer Interrupt.
[PID = 3] is running, variable: 15
[PID = 3] is running, variable: 16
[PID = 3] is running, variable: 17
```

switch to [PID = 2 COUNTER = 9]

```
[S] Supervisor Page Falut, scause: 00000000000000c, stval: 0000000000100e8, sepc: 0000000000100e8
pgd is 000000000000001f
[S] Supervisor Page Falut, scause: 00000000000000f, stval: 0000003fffffffff8, sepc: 000000000010124
pgd is 0000000000000017
[S] Supervisor Page Falut, scause: 00000000000000d, stval: 000000000011880, sepc: 000000000010140
pgd is 000000000000001f
[PID = 2] is running, variable: 0
[PID = 2] is running, variable: 1
[PID = 2] is running, variable: 2
[S] Supervisor Timer Interrupt.
[PID = 2] is running, variable: 3
[PID = 2] is running, variable: 4
```



```

[PID = 2] is running, variable: 4
[S] Supervisor Timer Interrupt.
[PID = 2] is running, variable: 5
[PID = 2] is running, variable: 6
[S] Supervisor Timer Interrupt.
[PID = 2] is running, variable: 7
[PID = 2] is running, variable: 8
[S] Supervisor Timer Interrupt.
[PID = 2] is running, variable: 9
[PID = 2] is running, variable: 10
[S] Supervisor Timer Interrupt.
[PID = 2] is running, variable: 11
[PID = 2] is running, variable: 12
[S] Supervisor Timer Interrupt.
[PID = 2] is running, variable: 13
[PID = 2] is running, variable: 14
[PID = 2] is running, variable: 15
[S] Supervisor Timer Interrupt.
[PID = 2] is running, variable: 16
[PID = 2] is running, variable: 17
[S] Supervisor Timer Interrupt.
[PID = 2] is running, variable: 18
[PID = 2] is running, variable: 19

switch to [PID = 1 COUNTER = 1]

[S] Supervisor Timer Interrupt.
[PID = 1] is running, variable: 9
[PID = 1] is running, variable: 10
[PID = 1] is running, variable: 11

switch to [PID = 2 COUNTER = 4]

[S] Supervisor Timer Interrupt.
[PID = 2] is running, variable: 20
[PID = 2] is running, variable: 21
[S] Supervisor Timer Interrupt.
[PID = 2] is running, variable: 22
[PID = 2] is running, variable: 23
[S] Supervisor Timer Interrupt.
[PID = 2] is running, variable: 24
[PID = 2] is running, variable: 25
[S] Supervisor Timer Interrupt.
[PID = 2] is running, variable: 26
[PID = 2] is running, variable: 27

switch to [PID = 4 COUNTER = 4]

```

```

switch to [PID = 4 COUNTER = 4]

[S] Supervisor Timer Interrupt.
[PID = 4] is running, variable: 11
[PID = 4] is running, variable: 12
[S] Supervisor Timer Interrupt.
[PID = 4] is running, variable: 13
[PID = 4] is running, variable: 14
[PID = 4] is running, variable: 15
[S] Supervisor Timer Interrupt.
[PID = 4] is running, variable: 16
[PID = 4] is running, variable: 17
[S] Supervisor Timer Interrupt.
[PID = 4] is running, variable: 18
[PID = 4] is running, variable: 19

switch to [PID = 3 COUNTER = 10]

[S] Supervisor Timer Interrupt.
[PID = 3] is running, variable: 18
[PID = 3] is running, variable: 19
[S] Supervisor Timer Interrupt.
[PID = 3] is running, variable: 20
[PID = 3] is running, variable: 21
[S] Supervisor Timer Interrupt.
[PID = 3] is running, variable: 22
[PID = 3] is running, variable: 23
[S] Supervisor Timer Interrupt.
[PID = 3] is running, variable: 24
[PID = 3] is running, variable: 25
[PID = 3] is running, variable: 26
[S] Supervisor Timer Interrupt.
[PID = 3] is running, variable: 27
[PID = 3] is running, variable: 28
[S] Supervisor Timer Interrupt.
[PID = 3] is running, variable: 29
[PID = 3] is running, variable: 30
[S] Supervisor Timer Interrupt.
[PID = 3] is running, variable: 31
[PID = 3] is running, variable: 32
[S] Supervisor Timer Interrupt.
[PID = 3] is running, variable: 33
[PID = 3] is running, variable: 34
[S] Supervisor Timer Interrupt.
[PID = 3] is running, variable: 35
[PID = 3] is running, variable: 36
[PID = 3] is running, variable: 37

```

## 思考题

1. `uint64_t vm_content_size_in_file`; 对应的文件内容的长度。为什么还需要这个域?

与lab4中的思考题相同, `vm_content_size_in_file`对应于文件中该段的大小, 而`vm_end - vm_start`是该段在内存中的大小。二者有不同的含义, 因为可加载段可能包含.bss节, 该节包含未初始化的数据。将此数据存储在磁盘或文件上会很浪费, 因此, 仅在ELF文件加载到内存后才占用空间。所以我们需要`vm_content_size_in_file`来标记对应的文件内容的长度, 这个长度是不能用`vm_end - vm_start`来表示的。

2. `struct vm_area_struct vmas[0]`; 为什么可以开大小为 0 的数组? 这个定义可以和前面的 `vma_cnt` 换个位置吗?

`struct vm_area_struct vmas[0]`; 定义柔性数组成员, 它允许在结构体的末尾定义一个数组, 但数组的大小可以是0,通常用于动态分配内存来保存可变大小的数据,因为每个进程都用了一个 PGSIZE 的大小来承载 `task_struct` 因此有足够的空间来让 `vmas[]` 动态的添加数据, 但 `vmas[]` 必须在最后, 不可以和 `vma_cnt` 换位置, 如果换了位置, `vmas[]`动态添加数据的时候会与`vma_cnt`产生内存冲突, 会覆盖掉 `vma_cnt`.